

# Design for Retry

**Microservices, REST, message busses  
and why idempotency is the only way to scale**

Aria Stewart

@aredridel

I'm here thanks to PayPal

I'm going to talk about errors.

It's going to be okay.

**We all know HTTP**

2xx OK

3xx Go elsewhere

4xx Tell user what they did wrong

5xx Bail out and log an error

**I'd call this**

**Error avoidance**

**You can't avoid errors**

**Here's the secret**

**Handle errors instead**

4xx Tell the user what they did wrong

5xx Save that request and do something with it later.

# Retry it

5xx are errors the *requestor* can handle



# Causes!

- database down
- bug in a service
- Deploy in progress
- power failure
- kicked a cable
- Network congestion
- Capacity exceeded

- Tree fell on the data center
- earthquake
- tornado
- birds, snakes and aeroplanes
- Black Friday
- Slashdot effect
- Interns
- QA tests
- DoS attack

**You need a queue**

# Lots of ways to do it

Database on the nodes

Log file

Queue server

# **gearman**

Queues built in

There are many alternatives, but gearmand is very simple.

The memcache of job queues.

# Three statuses:

- OK (Like 200)
- FAIL (Like 400)
- ERROR (Like 500)

**design so ERROR can  
be retried.**

gearmand **automatically tries a job**  
**ERROR again.**

And again.

And again.



# **If it isn't sure it worked?**

Tries it again.

**You *cannot* know if an  
error is a failure.**

# Error handling gets simpler

- Exception? ERROR.
- Database down? ERROR.
- Downstream service timeout? ERROR.

Maybe you retry right away.

**How many of you  
have used a job  
queue?**

**You have used a job  
queue**

# Let me tell you about one

TRILLIONS of messages

MILLIONS of nodes

100% availability (at least partial) for years.

32 years.

Resilient to MILLIONS of bad actors.

It is attached to the most malicious network.

# EMAIL.

250 OK

4xx RETRY

5xx Fail

# Responsibility for messages

250 - accept responsibility

4xx - reject responsibility

5xx - return responsibility



# **reject responsibility.**

If there's an error? Fail fast.

The requester can retry.

# Fail fast.

Queue work you can't reject. Reject everything you can if there is an error.

# You need a smart client.

Keeps outstanding requests.

Resubmit.

**We're Nodevember,  
We're Special  
Bonus content time!**

# **Smart Clients on the device**

**Toto, we're not in AWS anymore.**

**Ever lose an email  
because you've been  
logged out?**

**Ever try to write  
email on the web  
while not on the  
Internet?**

# Some ideas

- Queue things in `localStorage`
- Use third-party storage
- Integrate third-party services with this approach



**This is really good for  
offline-first design!**

**Being offline is the ultimate  
reliable error.**